```
import keras
keras.__version__
```

```
'2.4.3'
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
Mounted at /content/gdrive
```

```
from keras.applications import VGG16
```

```
conv_base = VGG16(weights='imagenet',
                  include_top=False,
                  input_shape=(150, 150, 3))
```

> Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16
> 58892288/58889256 [==============================] - 0s 0us/step

```
conv_base.summary()
```

```
Model: "vgg16"
_____
Layer (type)                 Output Shape              Param #
=================================================================
input_1 (InputLayer)         [(None, 150, 150, 3)]     0
_____
block1_conv1 (Conv2D)        (None, 150, 150, 64)      1792
_____
block1_conv2 (Conv2D)        (None, 150, 150, 64)      36928
_____
block1_pool (MaxPooling2D)   (None, 75, 75, 64)        0
_____
block2_conv1 (Conv2D)        (None, 75, 75, 128)       73856
_____
block2_conv2 (Conv2D)        (None, 75, 75, 128)       147584
_____
block2_pool (MaxPooling2D)   (None, 37, 37, 128)       0
_____
block3_conv1 (Conv2D)        (None, 37, 37, 256)       295168
_____
block3_conv2 (Conv2D)        (None, 37, 37, 256)       590080
_____
block3_conv3 (Conv2D)        (None, 37, 37, 256)       590080
_____
block3_pool (MaxPooling2D)   (None, 18, 18, 256)       0
_____
```

```
block4_conv1 (Conv2D)          (None, 18, 18, 512)        1180160
_____
block4_conv2 (Conv2D)          (None, 18, 18, 512)        2359808
_____
block4_conv3 (Conv2D)          (None, 18, 18, 512)        2359808
_____
block4_pool (MaxPooling2D)     (None, 9, 9, 512)          0
_____
block5_conv1 (Conv2D)          (None, 9, 9, 512)          2359808
_____
block5_conv2 (Conv2D)          (None, 9, 9, 512)          2359808
_____
block5_conv3 (Conv2D)          (None, 9, 9, 512)          2359808
_____
block5_pool (MaxPooling2D)     (None, 4, 4, 512)          0
=================================================================
Total params: 14,714,688
Trainable params: 14,714,688
Non-trainable params: 0
_____
```

```python
import os
import numpy as np
from keras.preprocessing.image import ImageDataGenerator

base_dir = '/content/gdrive/MyDrive/ML Assignment 2/cats_and_dogs_2'

train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validation')
test_dir = os.path.join(base_dir, 'test')

datagen = ImageDataGenerator(rescale=1./255)
batch_size = 20

def extract_features(directory, sample_count):
    features = np.zeros(shape=(sample_count, 4, 4, 512))
    labels = np.zeros(shape=(sample_count))
    generator = datagen.flow_from_directory(
        directory,
        target_size=(150, 150),
        batch_size=batch_size,
        class_mode='binary')
    i = 0
    for inputs_batch, labels_batch in generator:
        features_batch = conv_base.predict(inputs_batch)
        features[i * batch_size : (i + 1) * batch_size] = features_batch
        labels[i * batch_size : (i + 1) * batch_size] = labels_batch
        i += 1
        if i * batch_size >= sample_count:
            # Note that since generators yield data indefinitely in a loop,
            # we must `break` after every image has been seen once.
            break
    return features, labels
```

```
train_features, train_labels = extract_features(train_dir, 2000)
validation_features, validation_labels = extract_features(validation_dir, 1000)
test_features, test_labels = extract_features(test_dir, 1000)
```

```
    Found 3010 images belonging to 2 classes.
    ---------------------------------------------------------------------------
    KeyboardInterrupt                         Traceback (most recent call last)
    <ipython-input-6-e489372cb3cc> in <module>()
    ----> 1 train_features, train_labels = extract_features(train_dir, 2000)
          2 validation_features, validation_labels = extract_features(validation_dir, 1000)
          3 test_features, test_labels = extract_features(test_dir, 1000)

                              ⌃⌄ 4 frames

    /usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/utils.py in load_img(pa
    color_mode, target_size, interpolation)
        112                                 'The use of `load_img` requires PIL.')
        113         with open(path, 'rb') as f:
    --> 114             img = pil_image.open(io.BytesIO(f.read()))
        115             if color_mode == 'grayscale':
        116                 # if image is not already an 8-bit, 16-bit or 32-bit grayscale image

    KeyboardInterrupt:
```

```
train_features = np.reshape(train_features, (2000, 4 * 4 * 512))
validation_features = np.reshape(validation_features, (1000, 4 * 4 * 512))
test_features = np.reshape(test_features, (1000, 4 * 4 * 512))
```

```
    ---------------------------------------------------------------------------
    NameError                                 Traceback (most recent call last)
    <ipython-input-17-242630fb165b> in <module>()
          1
    ----> 2 train_features = np.reshape(train_features, (2000, 4 * 4 * 512))
          3 validation_features = np.reshape(validation_features, (1000, 4 * 4 * 512))
          4 test_features = np.reshape(test_features, (1000, 4 * 4 * 512))

    NameError: name 'train_features' is not defined
```

```
from keras import models
from keras import layers
from keras import optimizers

model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
```

```python
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(train_features, train_labels,
                    epochs=30,
                    batch_size=20,
                    validation_data=(validation_features, validation_labels))


import matplotlib.pyplot as plt

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()


from keras import models
from keras import layers

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))


model.summary()


print('This is the number of trainable weights '
      'before freezing the conv base:', len(model.trainable_weights))
```

```python
conv_base.trainable = False


print('This is the number of trainable weights '
      'after freezing the conv base:', len(model.trainable_weights))


from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
      rescale=1./255,
      rotation_range=40,
      width_shift_range=0.2,
      height_shift_range=0.2,
      shear_range=0.2,
      zoom_range=0.2,
      horizontal_flip=True,
      fill_mode='nearest')

# Note that the validation data should not be augmented!
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
        # This is the target directory
        train_dir,
        # All images will be resized to 150x150
        target_size=(150, 150),
        batch_size=20,
        # Since we use binary_crossentropy loss, we need binary labels
        class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
        validation_dir,
        target_size=(150, 150),
        batch_size=20,
        class_mode='binary')

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=2e-5),
              metrics=['acc'])

history = model.fit_generator(
      train_generator,
      steps_per_epoch=100,
      epochs=30,
      validation_data=validation_generator,
      validation_steps=50,
      verbose=2)


model.save('cats_and_dogs_small_3.h5')
```

```python
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()


conv_base.summary()


conv_base.trainable = True

set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False


model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-5),
              metrics=['acc'])

history = model.fit_generator(
      train_generator,
      steps_per_epoch=100,
      epochs=100,
      validation_data=validation_generator,
      validation_steps=50)


model.save('cats_and_dogs_small_4.h5')
```

```
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()


def smooth_curve(points, factor=0.8):
  smoothed_points = []
  for point in points:
    if smoothed_points:
      previous = smoothed_points[-1]
      smoothed_points.append(previous * factor + point * (1 - factor))
    else:
      smoothed_points.append(point)
  return smoothed_points

plt.plot(epochs,
         smooth_curve(acc), 'bo', label='Smoothed training acc')
plt.plot(epochs,
         smooth_curve(val_acc), 'b', label='Smoothed validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs,
         smooth_curve(loss), 'bo', label='Smoothed training loss')
plt.plot(epochs,
         smooth_curve(val_loss), 'b', label='Smoothed validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()


test_generator = test_datagen.flow_from_directory(
        test_dir
```

```
        test_dir,
        target_size=(150, 150),
        batch_size=20,
        class_mode='binary')

test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)
print('test acc:', test_acc)
```